


Multicore Systems

Sujoy Sinha Roy
School of Computer Science
University of Birmingham


Recap: Simple von Neumann Computer



We started C programming assuming a simple von Neumann computation model.

- We executed program on a single CPU (or processor) step-by-step.
- We stored data and program instructions in the Memory.

Recap: Memory Hierarchy



- We studied memory hierarchy from a programmer's perspective
- Learnt how to write efficient C code to exploit memory hierarchy
 - Programs are still executed on a single CPU (or processor) step-by-step but more efficiently



Produced	From May 7, 1997 to December 26, 2003 ^[1]
Common manufacturer(s)	Intel
Max. CPU clock rate	233 MHz to 450 MHz
FSB speeds	66 MHz to 100 MHz
Min. feature size	0.35 μm to 0.18 μm
Instruction set	IA-32, MMX



Produced	From February 26, 1999 to May 18, 2007 ^[1]
Common manufacturer(s)	Intel
Max. CPU clock rate	450 MHz to 1.4 GHz
FSB speeds	100 MHz to 133 MHz
Min. feature size	0.25 μm to 0.13 μm
Instruction set	IA-32, MMX, SSE

Clock frequency almost doubled w.r.t Pentium II



Produced	From November 20, 2000 to August 8, 2008
Max. CPU clock rate	1.3 GHz to 3.8 GHz
FSB speeds	400 MT/s to 1066 MT/s
Instruction set	x86 (i386), x86-64 (only some chips), MMX, SSE, SSE2, SSE3

Clock frequency almost doubled w.r.t Pentium III


What is the clock frequency in 2020?




Intel® Core™ i7-10510U Processor (8M Cache, Up to 4.80 GHz)

- 8 MB Intel® Smart Cache
- 4 Cores
- 8 Threads
- 4.90 GHz Max Turbo Frequency
- U - Ultra-low power
- 10th Generation

Little increase in clock frequency after 10 years since Pentium 4



Source: Herb Sutter "The Free Lunch Is Over: A Fundamental Turn Toward Concurrency in Software" available at <http://www.gotw.ca/publications/concurrency-ddj.htm>



Source: Patrick Gelsinger, Intel Developer's Forum, Intel Corporation, 2004.

Applications are demanding more resources!

Alternative solution:

- Put many processing cores on the microprocessor chip.
- The number of cores doubles with each generation.

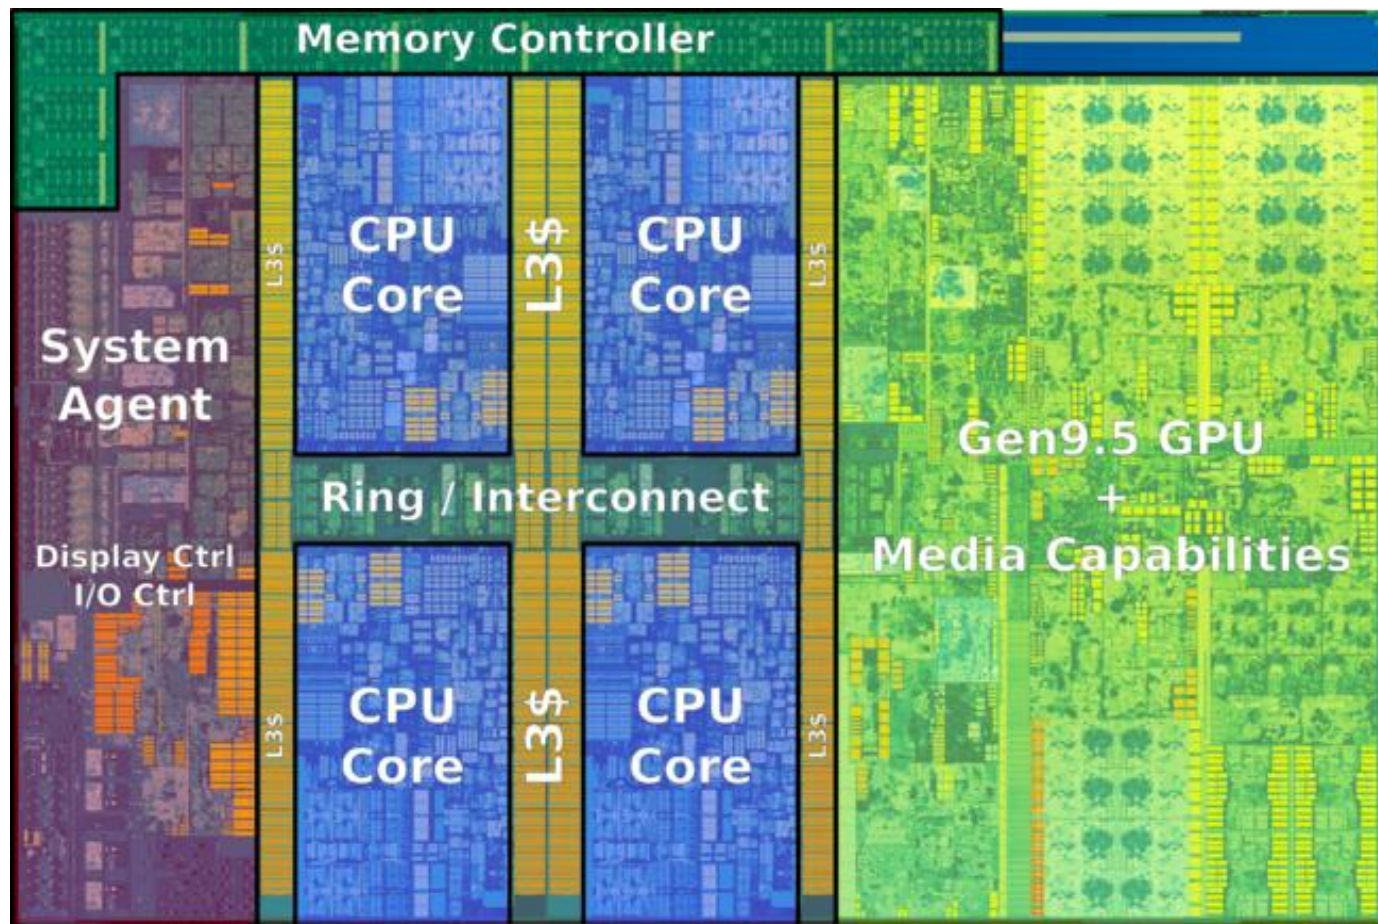


Photo of the quad-core GT2 Kaby Lake processor chip.
Used in mainstream desktop computers.

It has 4 CPU cores to compute in parallel.



Samsung Galaxy S10+

Exynos 9 Series (9820)

6.4" Quad HD+ Curved Dynamic AMOLED (3040x1440)


Android 9 (Pie)

8/12GB RAM, 128/256GB/1TB Storage, microSD(up to 512GB)


Rear: 12MP AF(Wide)+12MP AF(Tele)+16MP FF(Ultra Wide),
Front: 10MP AF+8MP FF(RGB Depth)

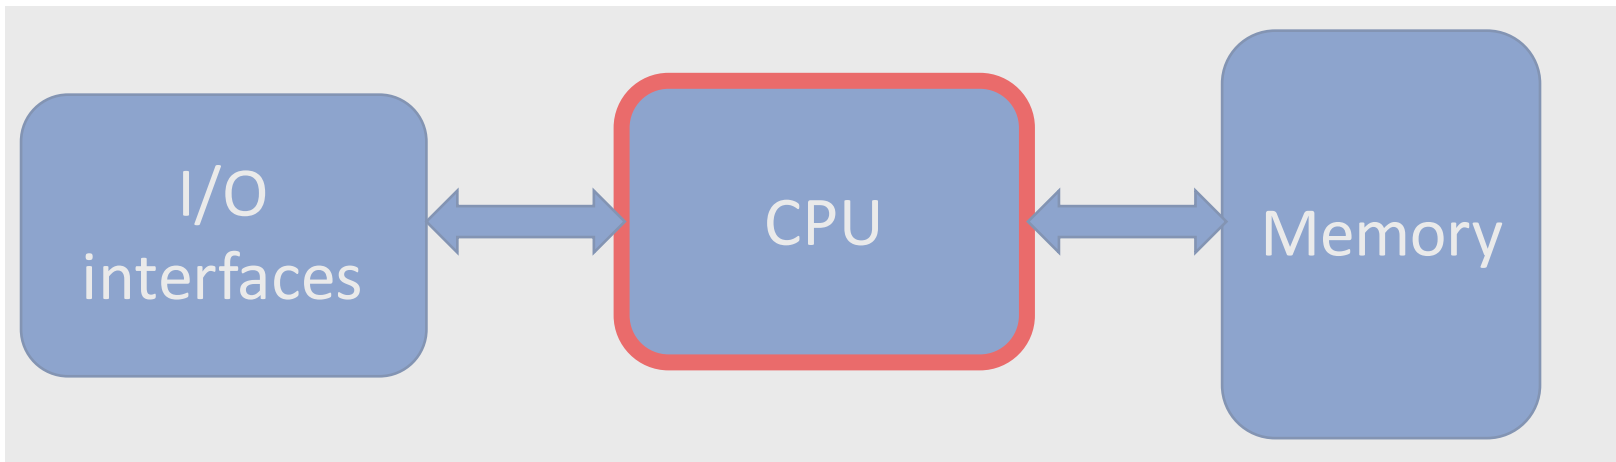
Superior Performance for Seamless Multi-tasking

Featuring a 4th generation custom CPU, the Exynos 9820's innovative tri-cluster architecture delivers premium processing power. The CPU consists of two custom cores for ultimate processing power, two Cortex-A75 cores for optimal performance, and four Cortex-A55 cores for greater efficiency, resulting in superior performance that lasts. Tri-cluster with intelligent task scheduler boosts multi-core performance by 15 percent when compared to the Exynos 9810, while the 4th generation custom CPU with enhanced memory access capability and cutting-edge architecture design improves single core performance by up to 20 percent or boosts power efficiency by up to 40 percent.¹




From Single Core
To
Multi Core



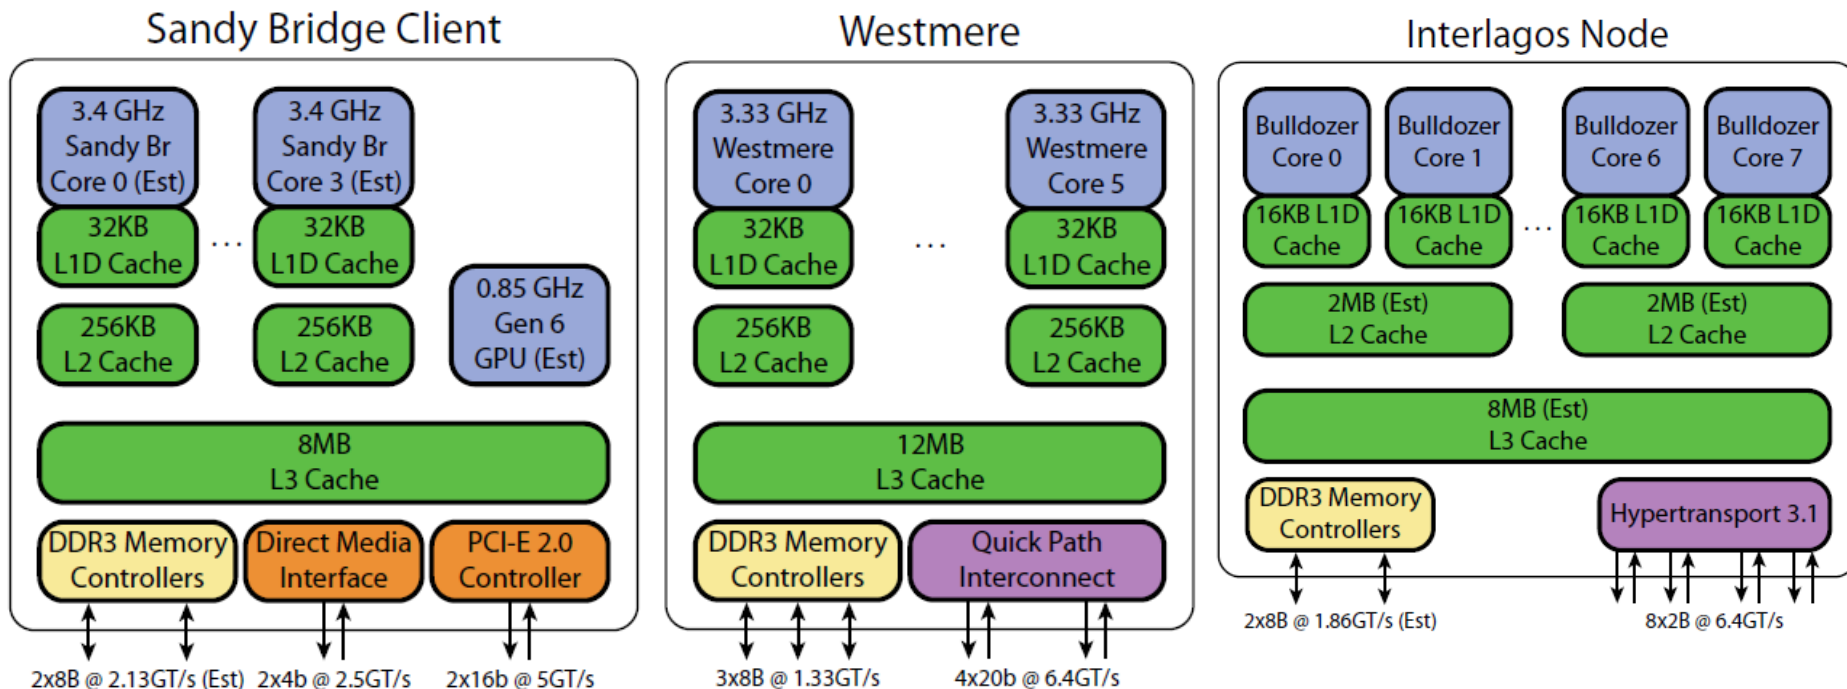


From Single Core
To
Multi Core

What happens to the Memory?
One memory? Multiple memory?

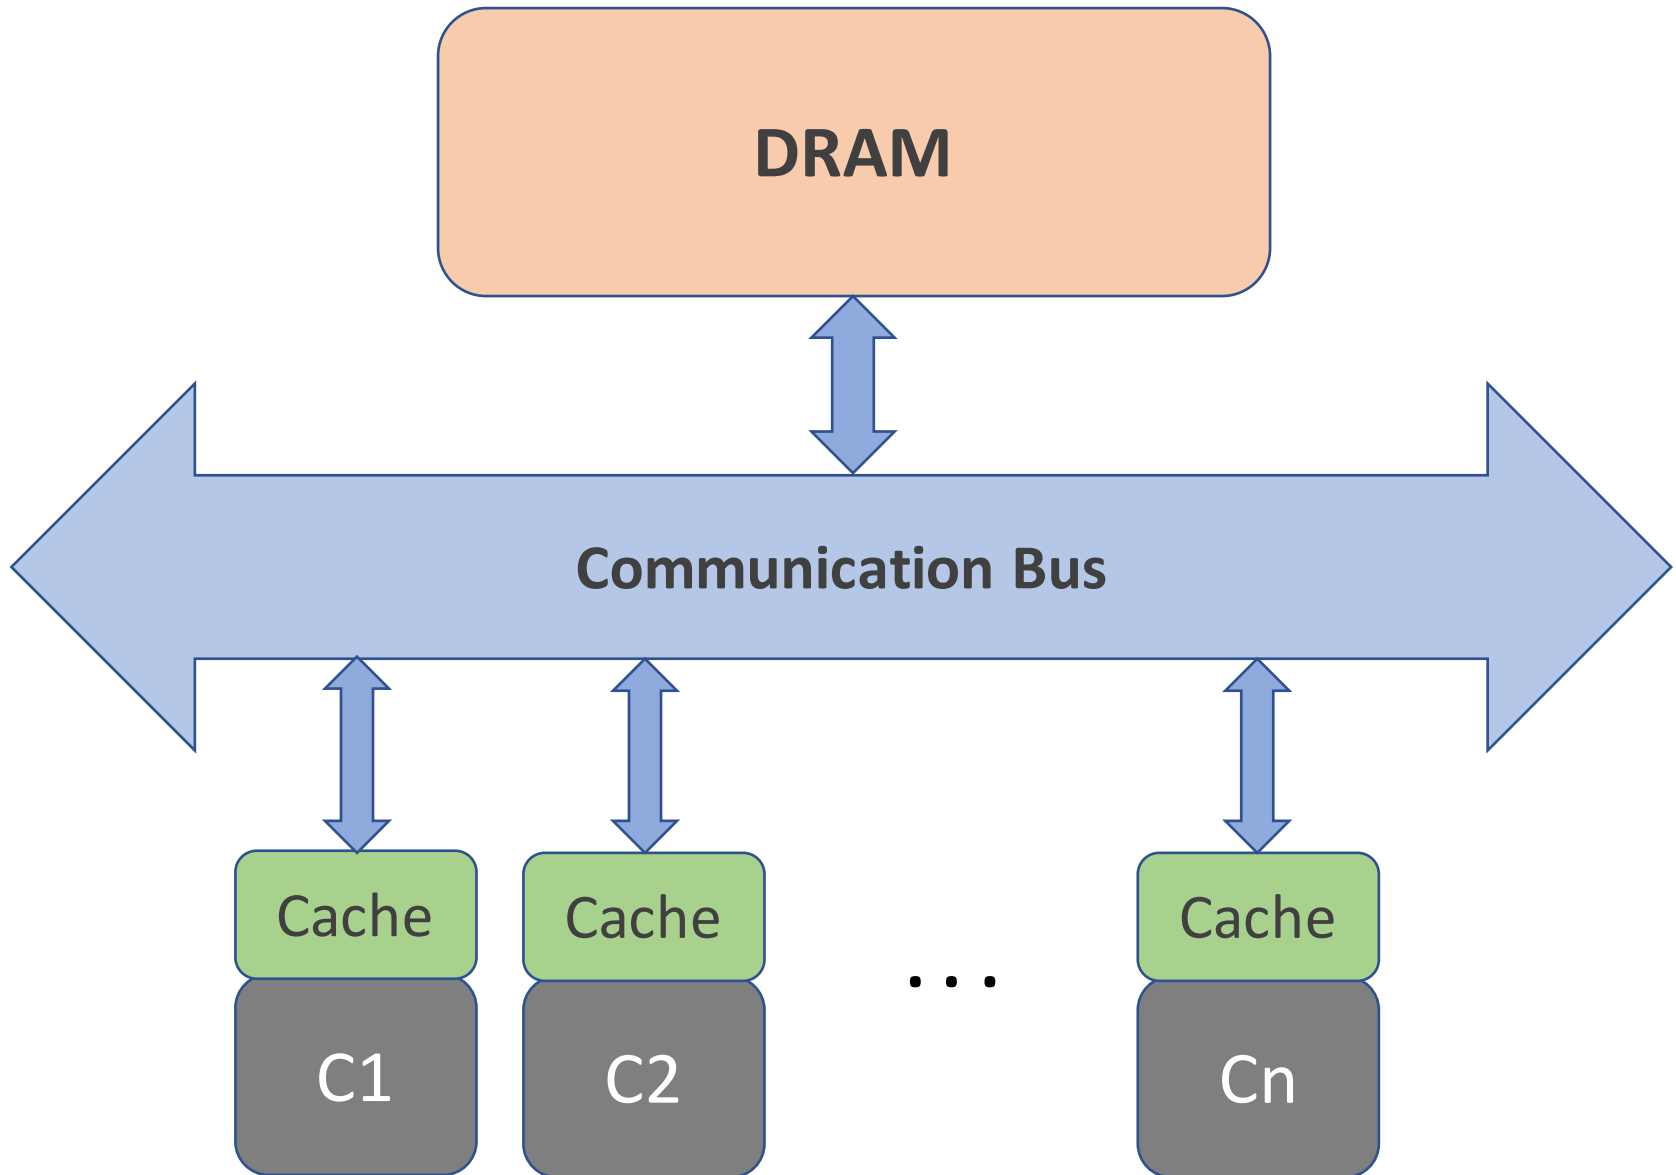


Examples of Intel Multicore Processors




- Each Core has its own L1 Cache
- In some processors, each core has its own L2 Cache
- All cores share one L3 cache
- All cores share one DDR3 DRAM (main memory)

Simplified Multicore Architecture



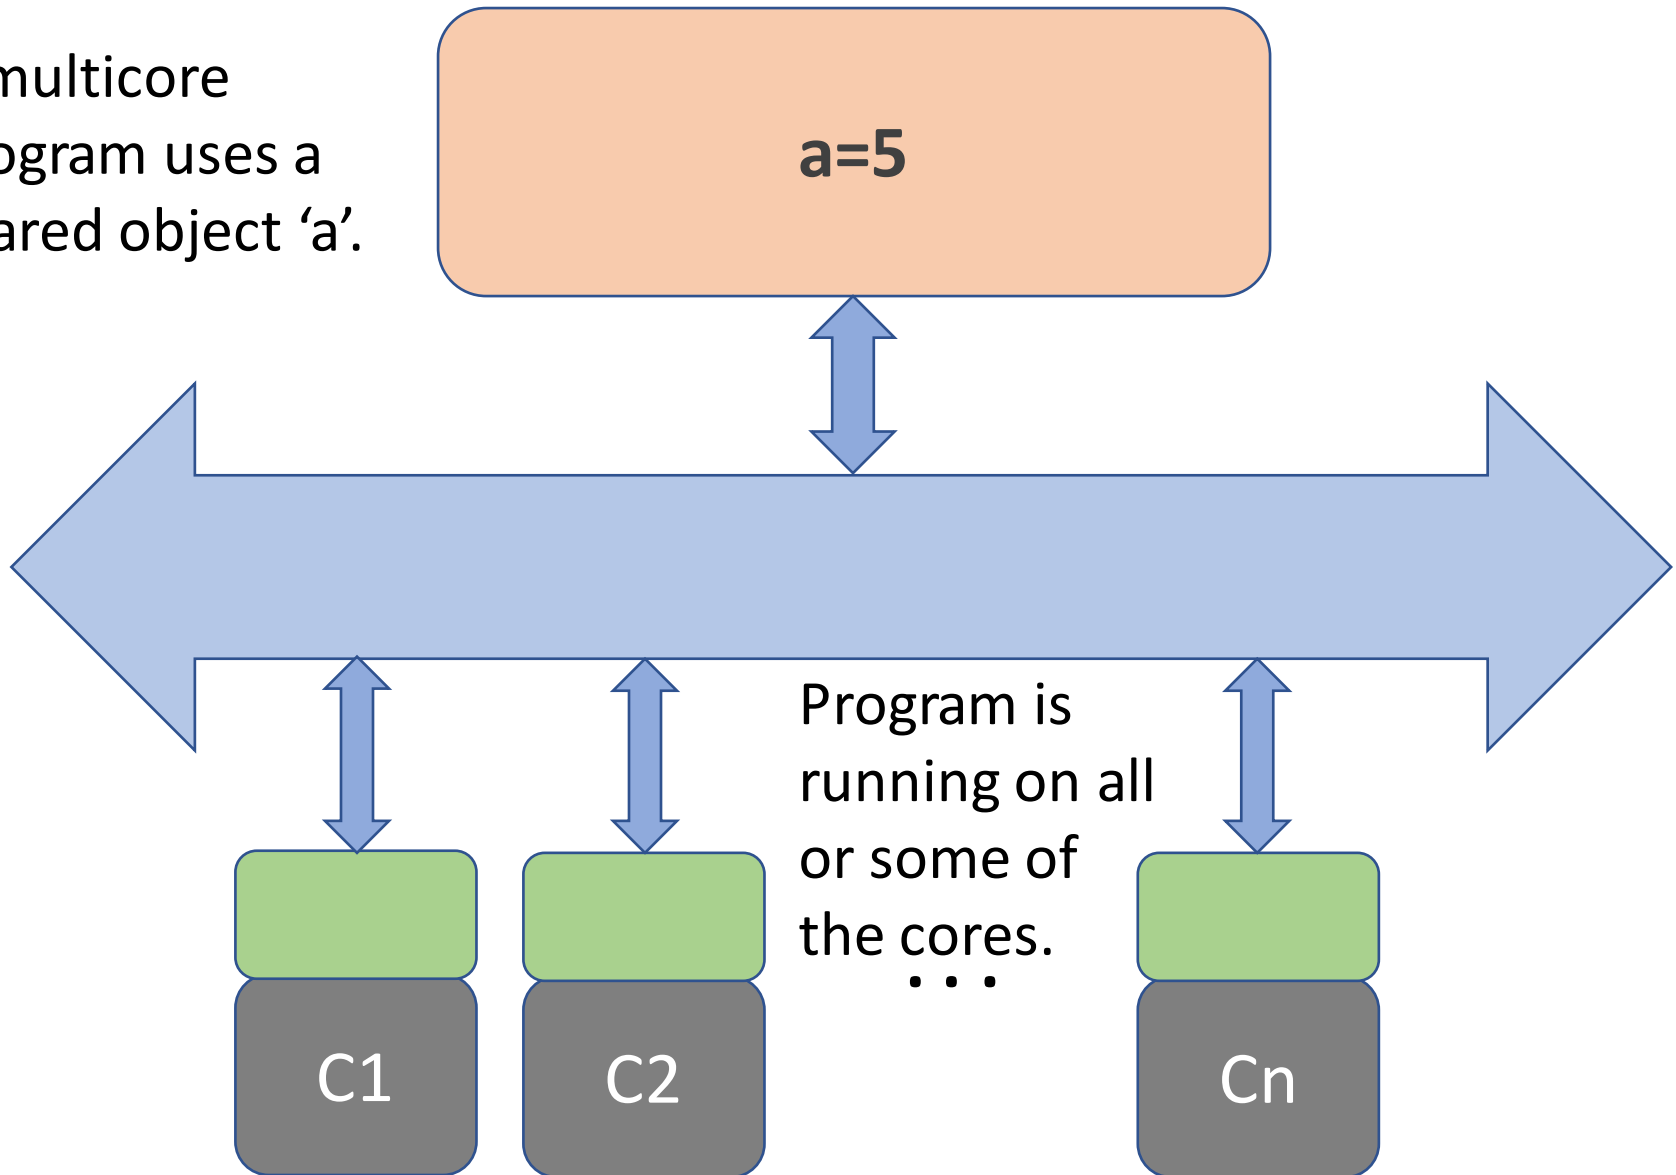
Coherence of data




- 'Coherence' is the quality of being logical and consistent. During program execution, data must remain coherent.
- On a multicore system, a data variable may reside in multiple caches and might get updated 'locally' by its local core
→ this causes coherence problem

Example: Cache coherence problem


A multicore program uses a shared object 'a'.



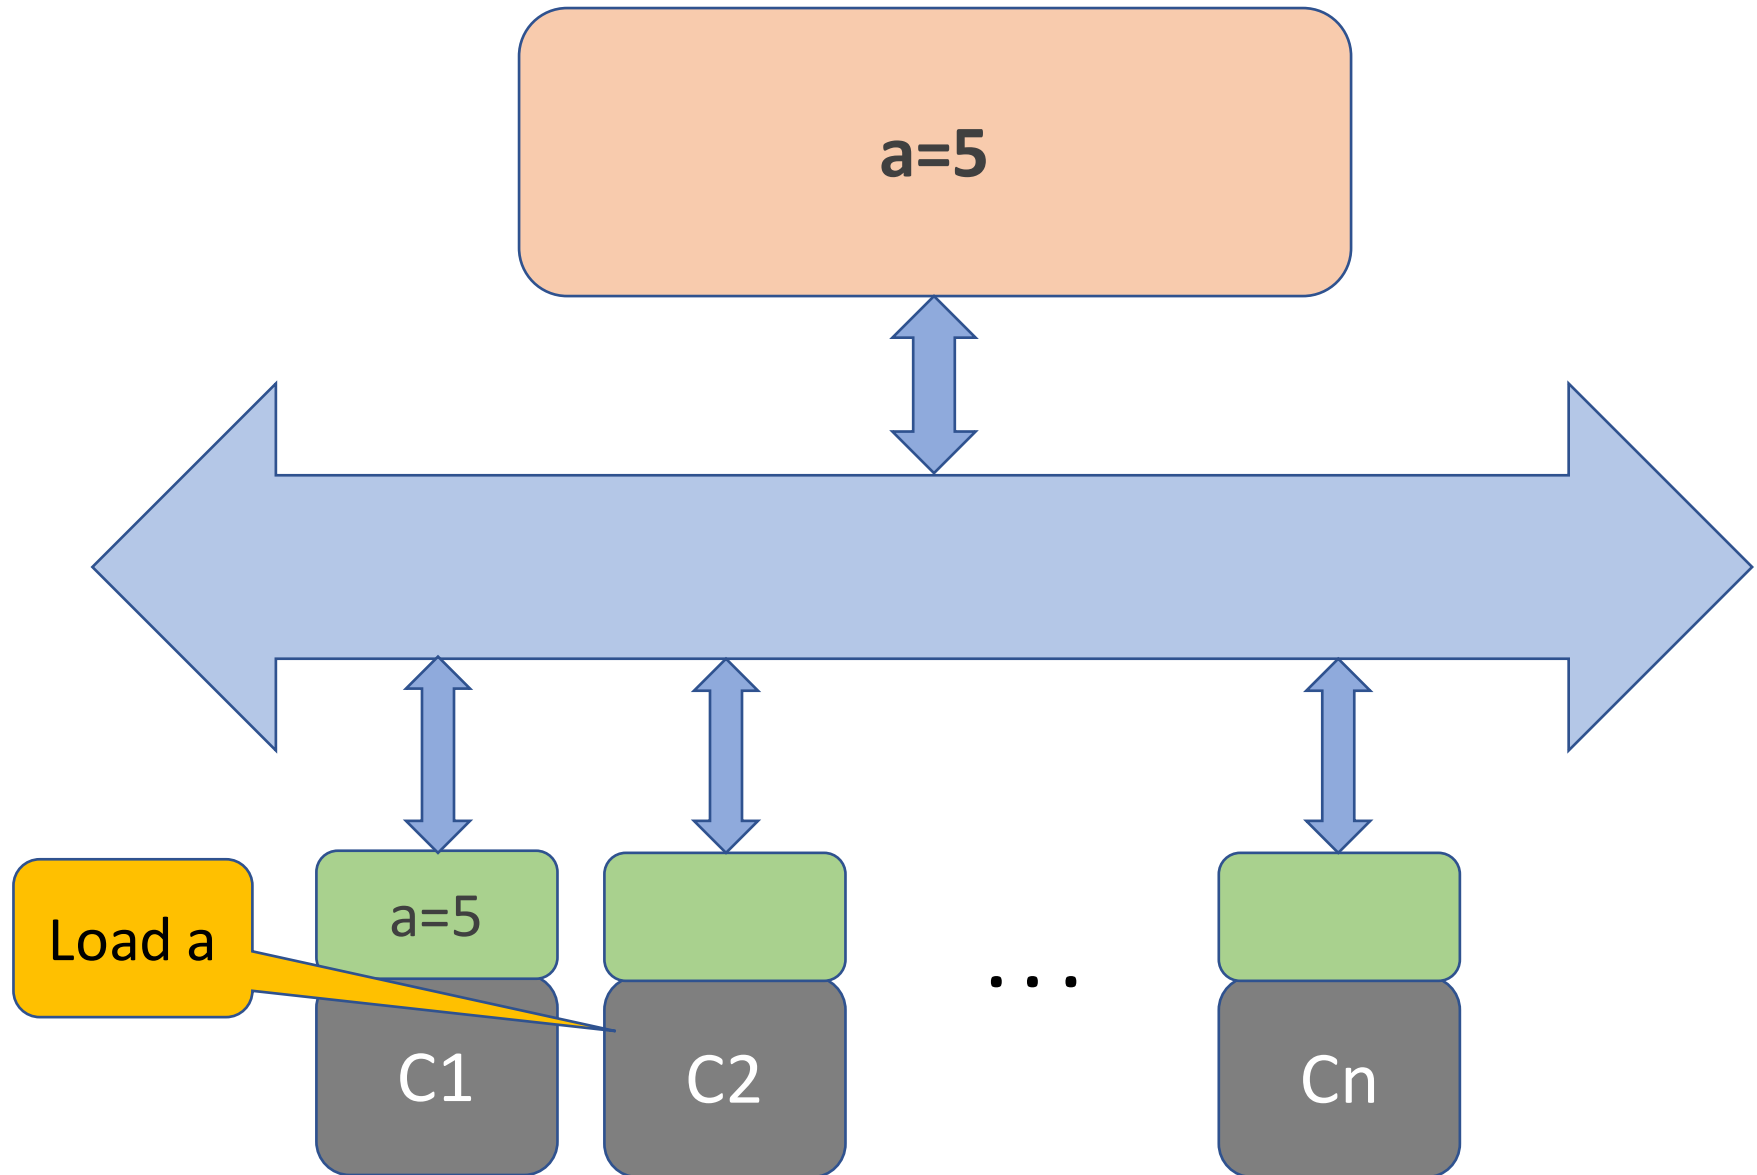
Example: Cache coherence problem



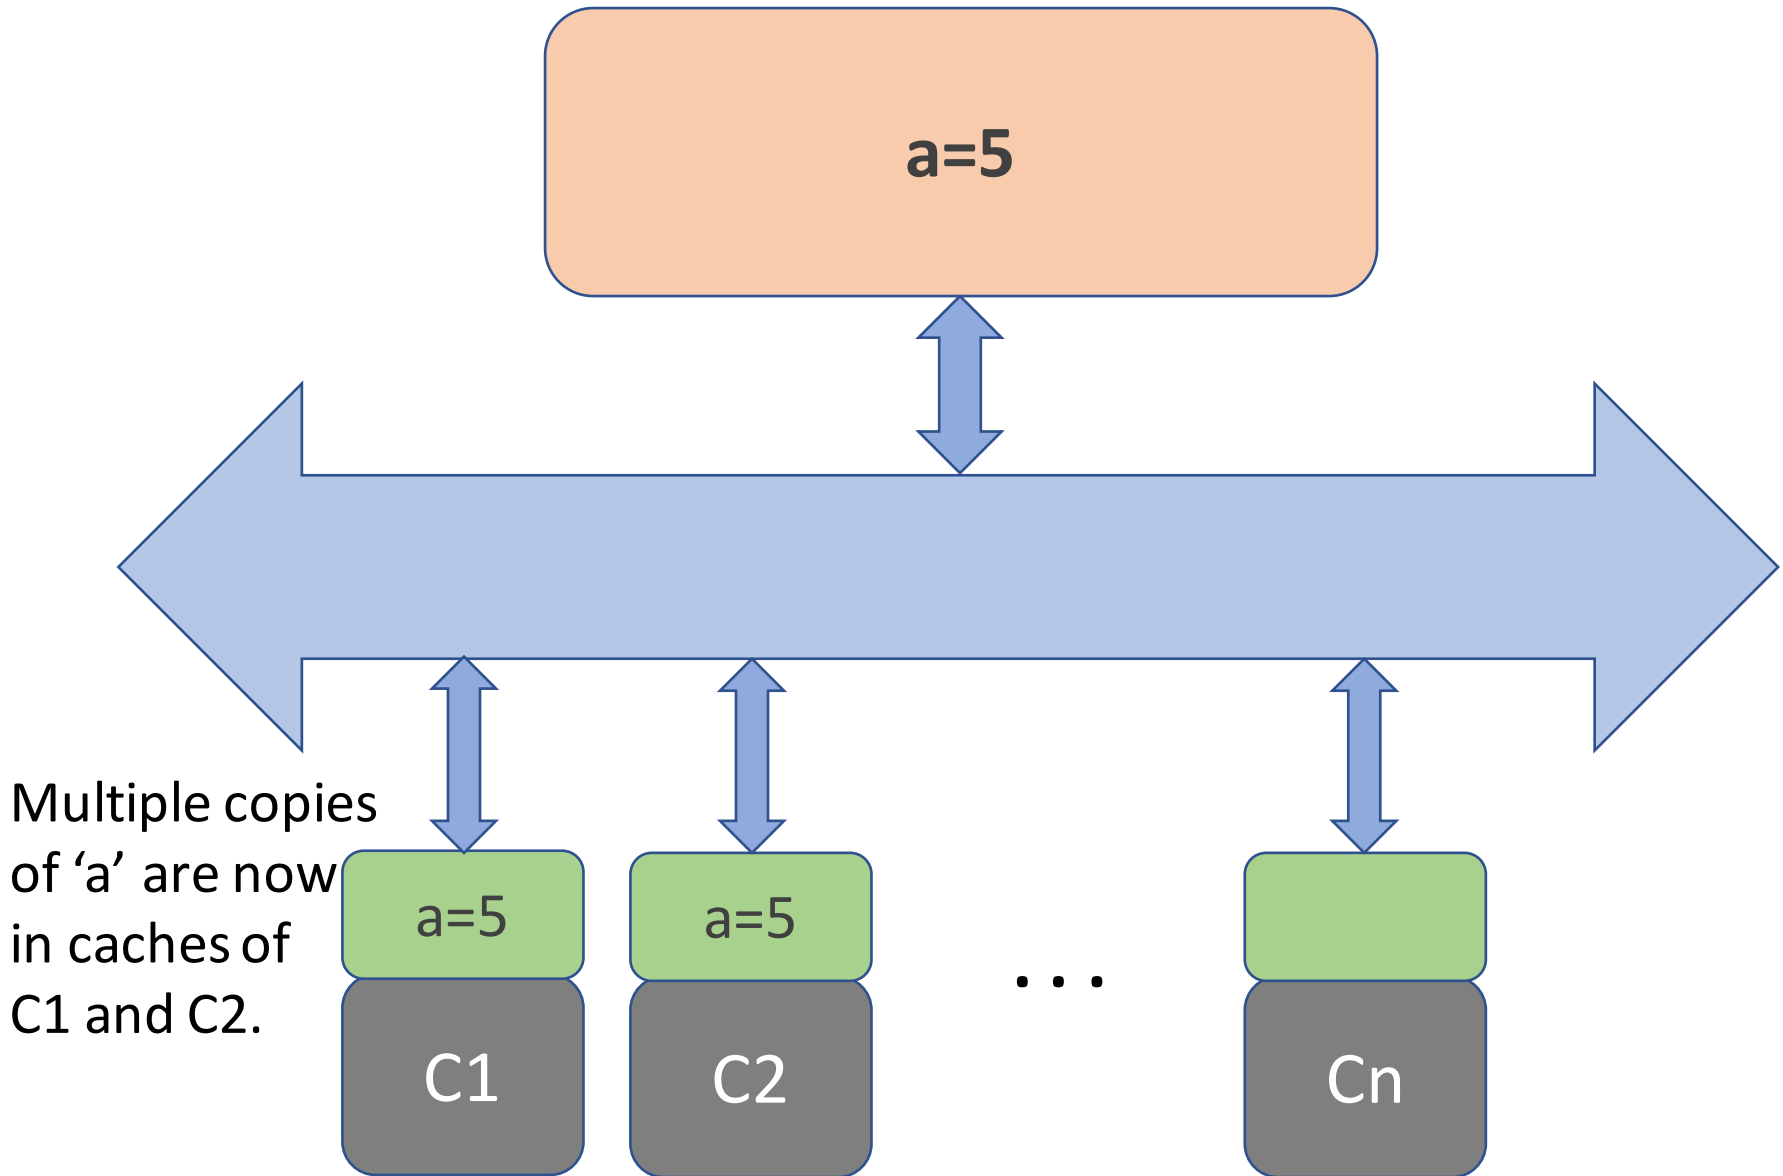
Example: Cache coherence problem



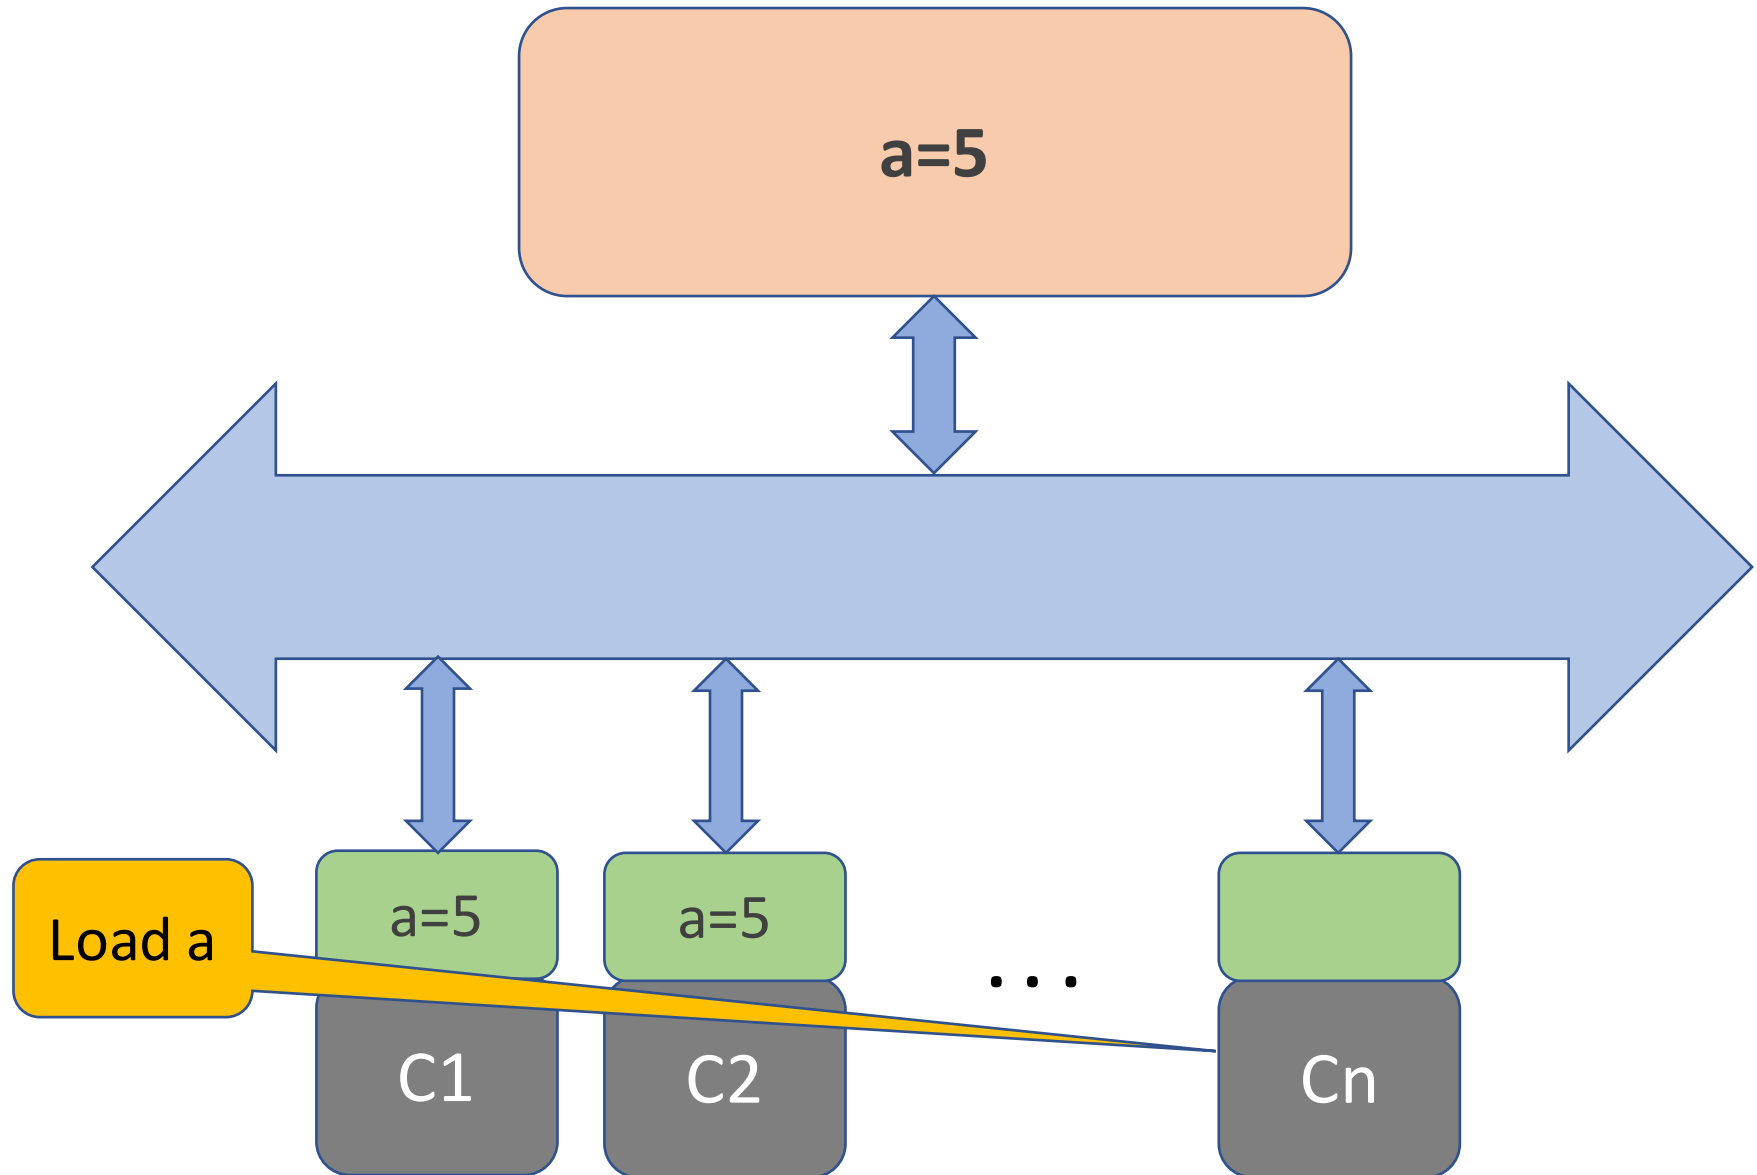
Example: Cache coherence problem



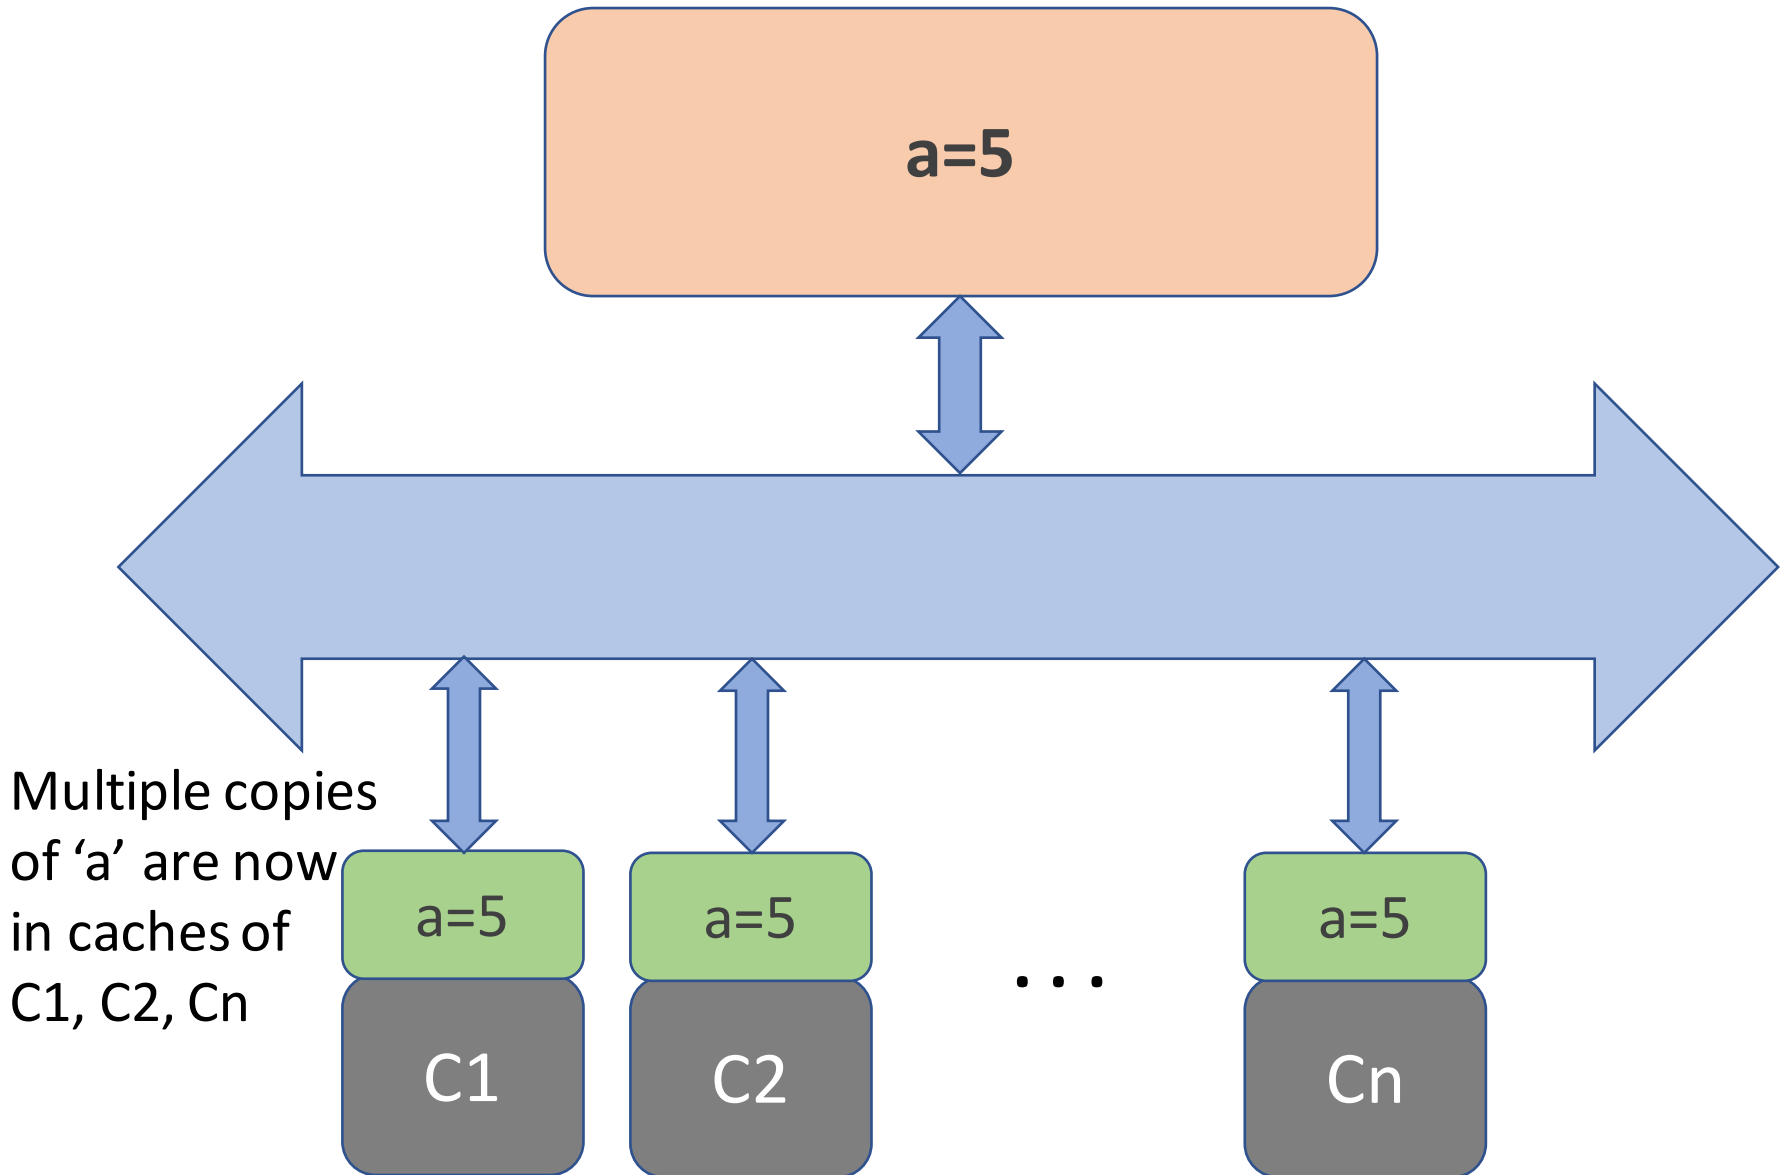
Example: Cache coherence problem



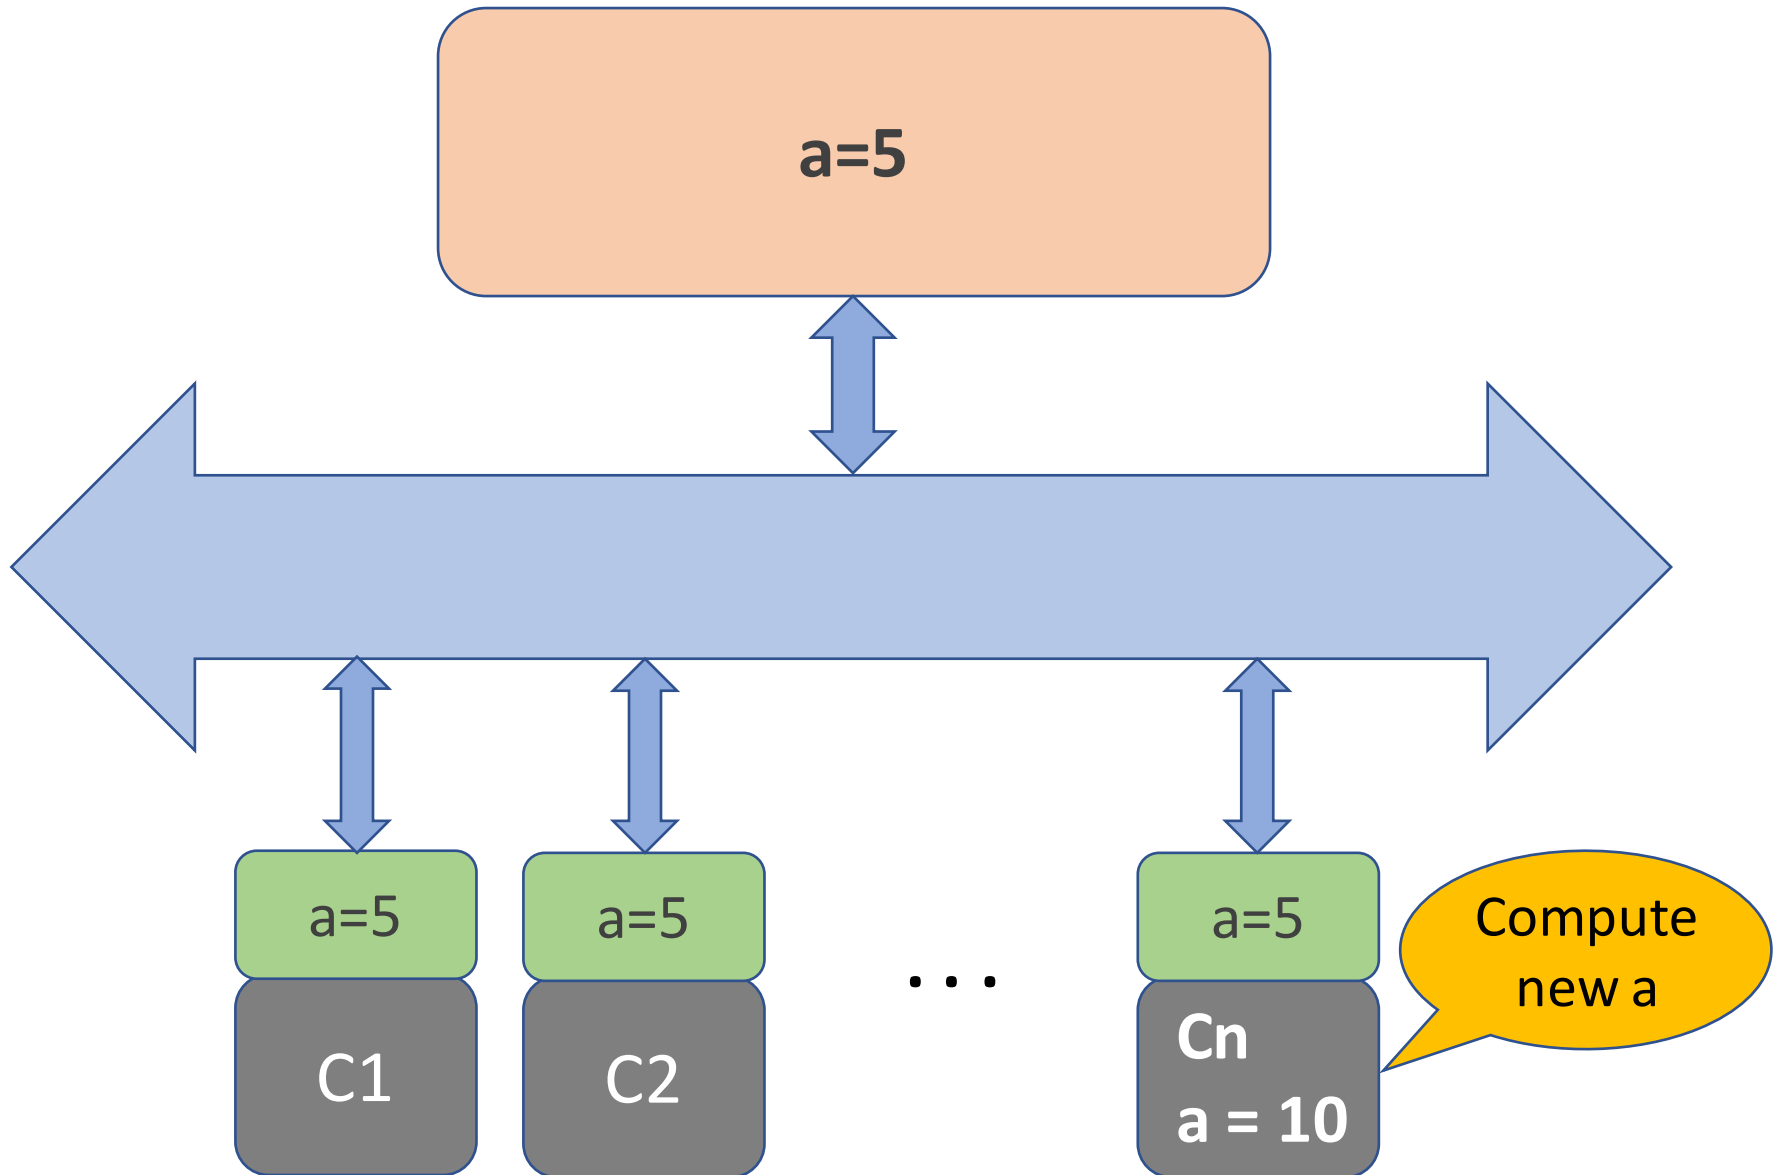
Example: Cache coherence problem



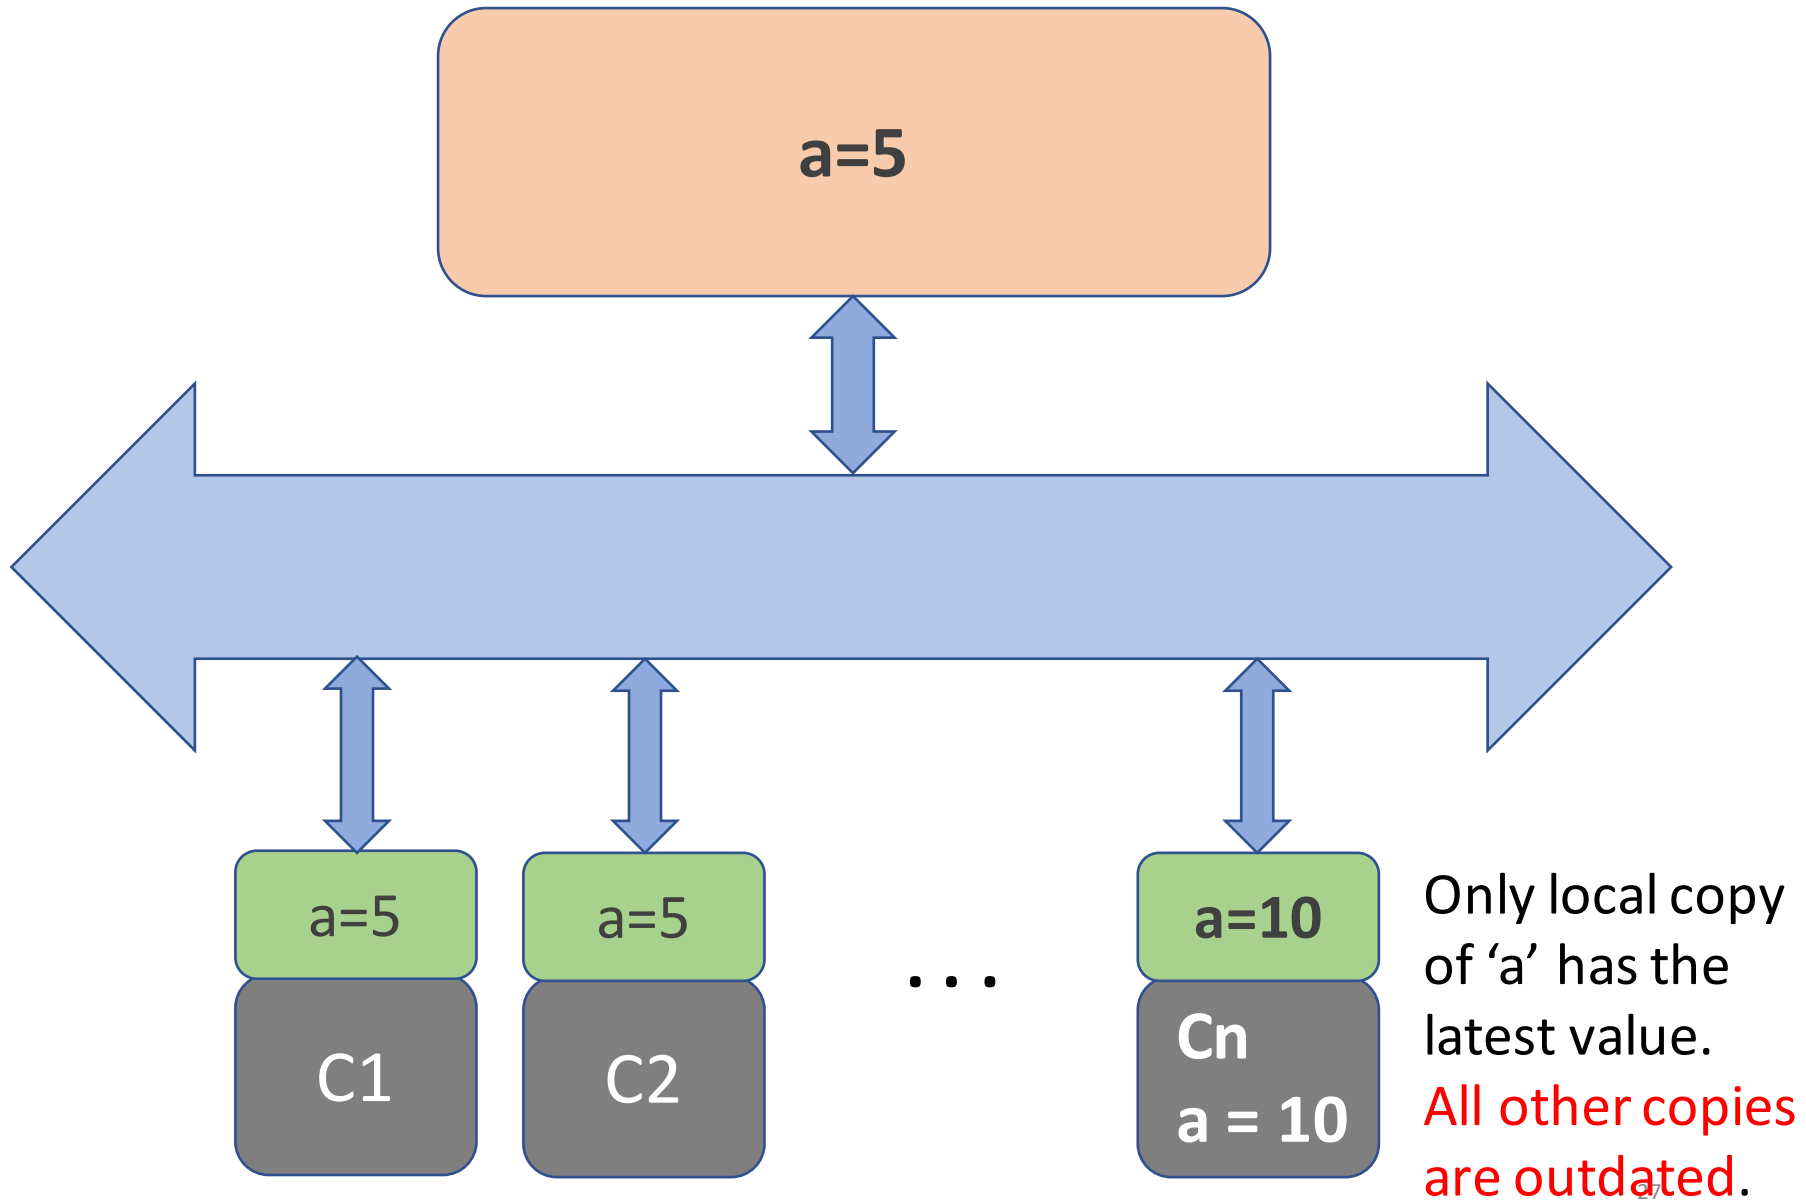
Example: Cache coherence problem



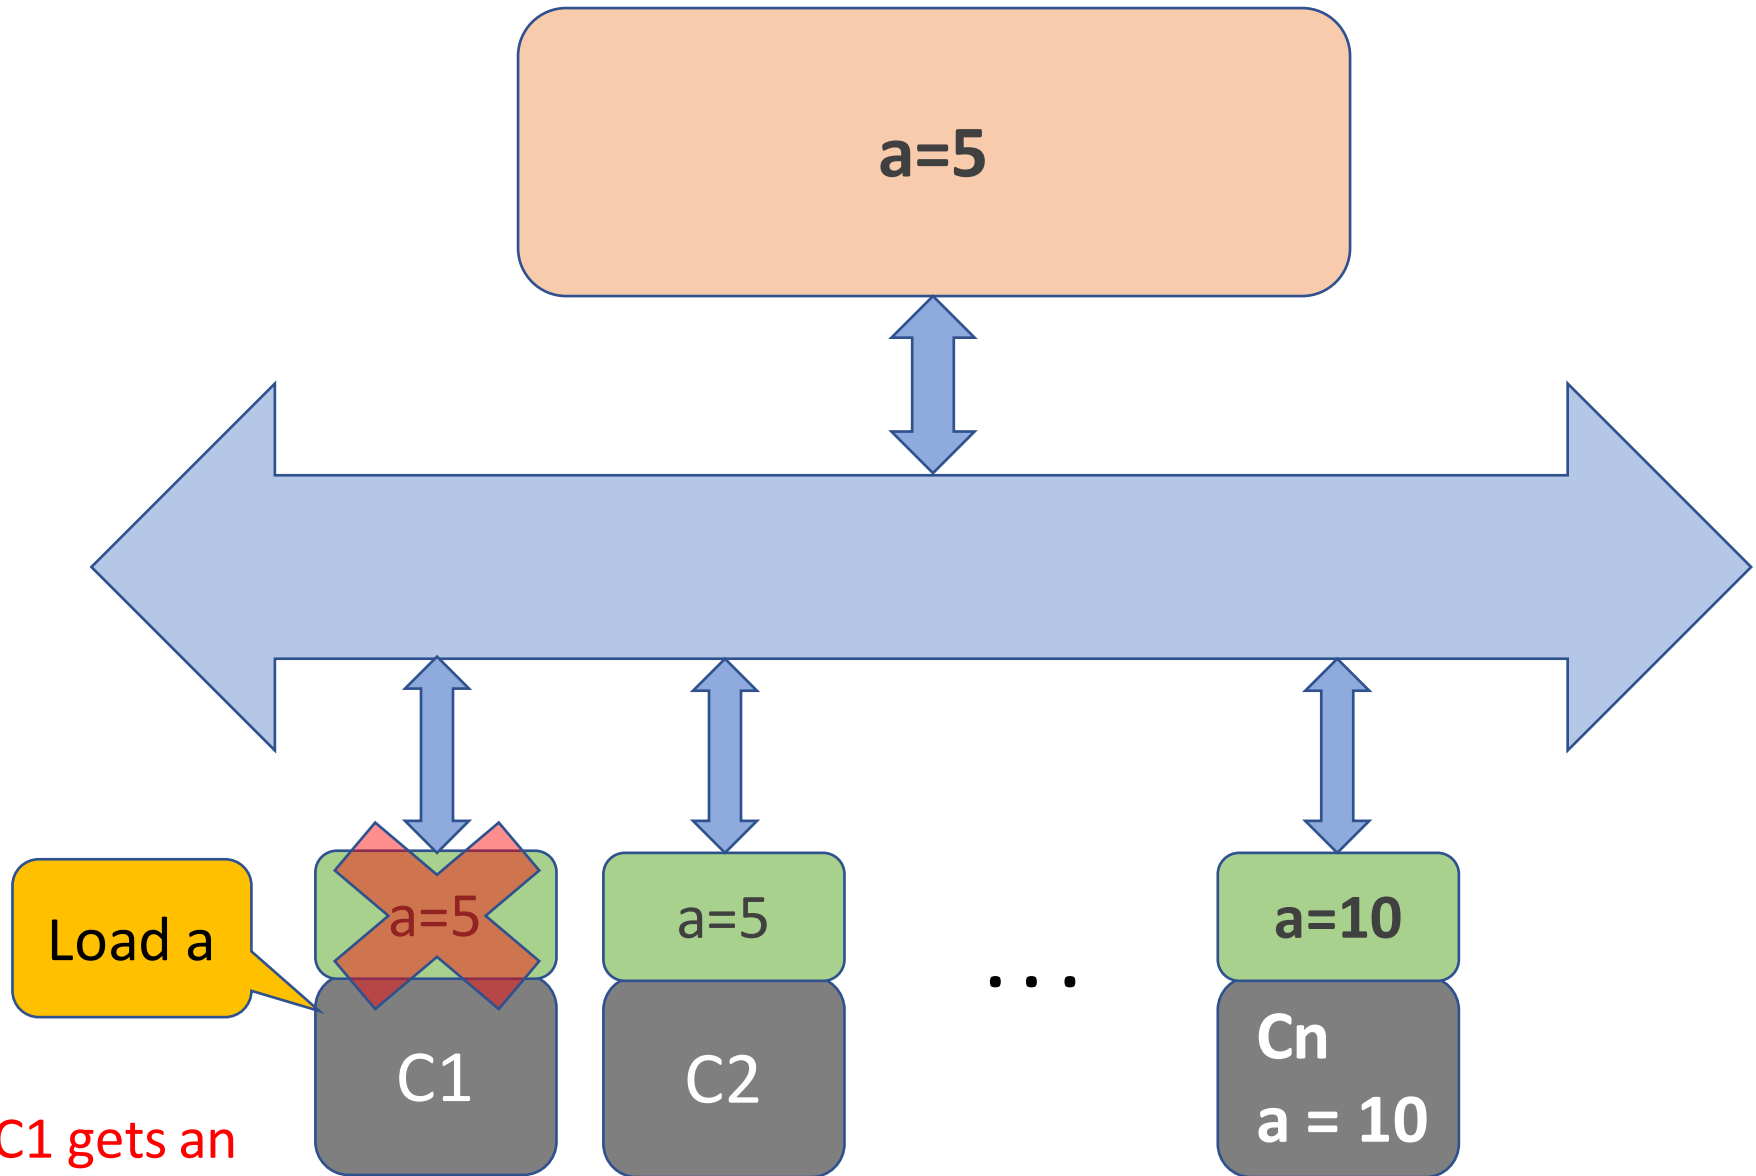
Example: Cache coherence problem



Example: Cache coherence problem



Example: Cache coherence problem

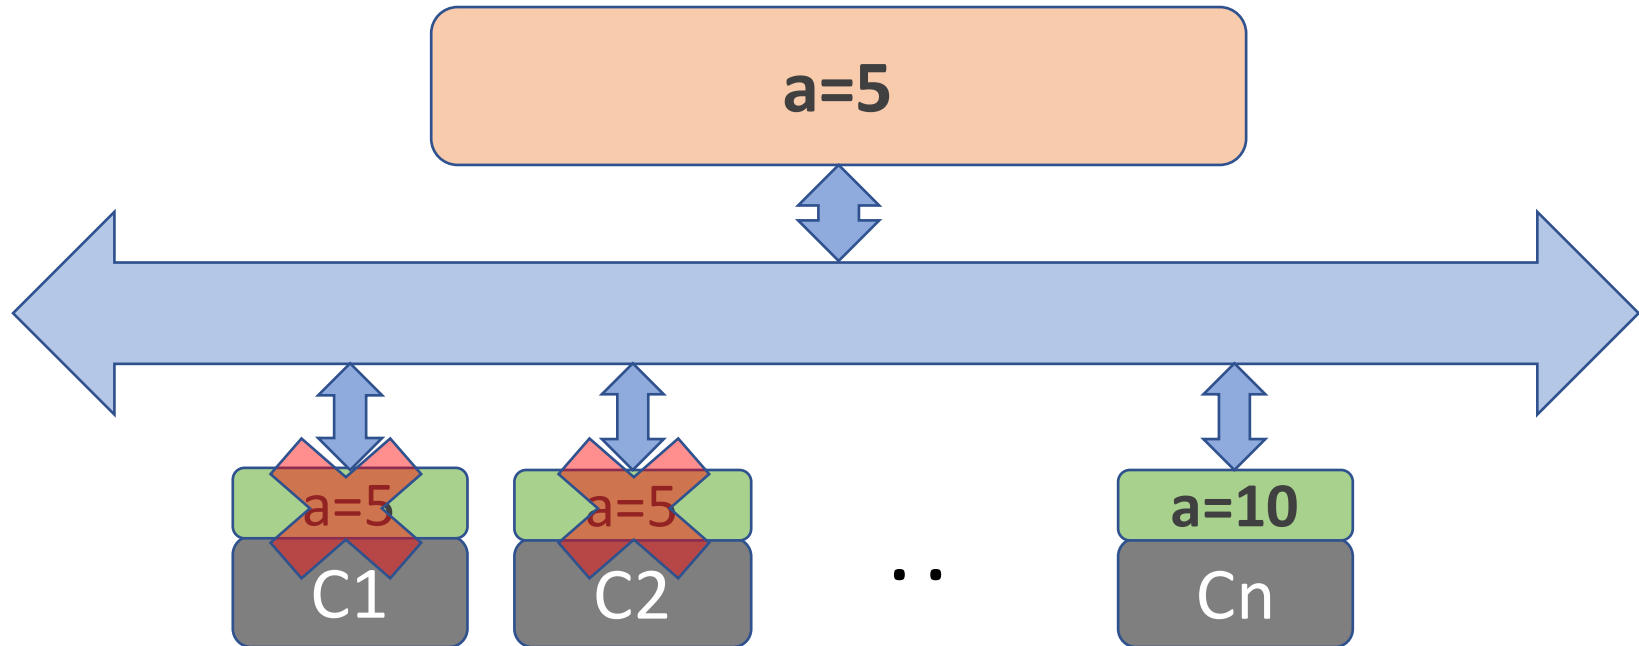


C1 gets an
outdated 'a' from its Cache.

Coherence mechanisms

To achieve coherence:

a **write to a memory location** must cause all other copies of this location to be **removed from the caches** they are in.



Do: As soon as 'a' gets updated by a core, invalidate or delete all other copies of it.

Coherence protocols

- Coherence protocols apply cache coherence in multicore systems.
- Goal is that two cores **must never see different values** for the same shared data.

How this is achieved?

Additional information is stored and passed

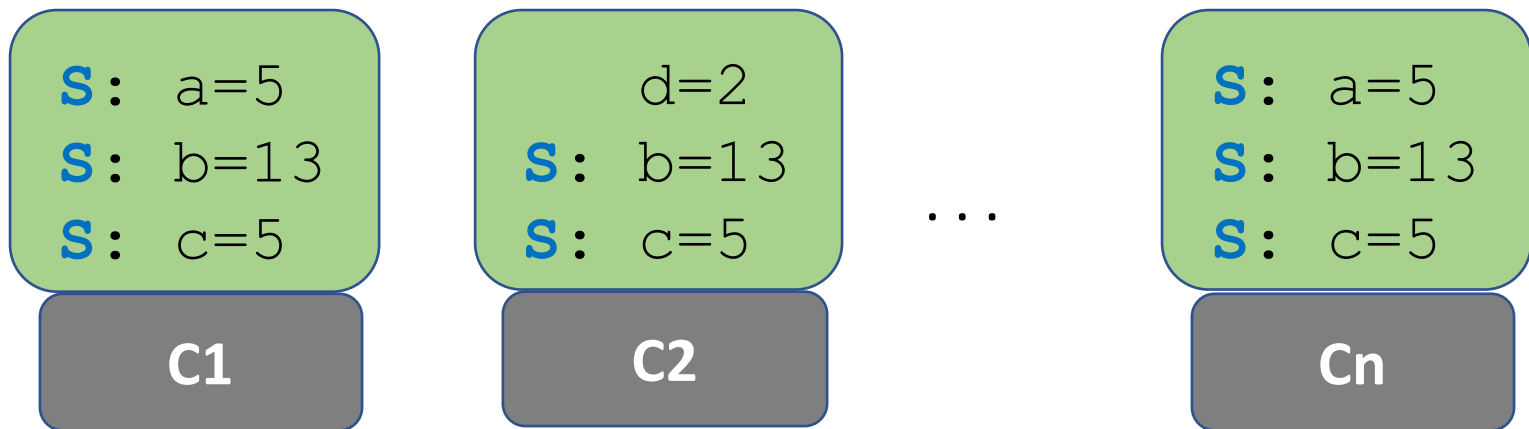
- Is this data stored in multiple caches?
- Has this cache line been modified?

MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Initial state of caches before computation

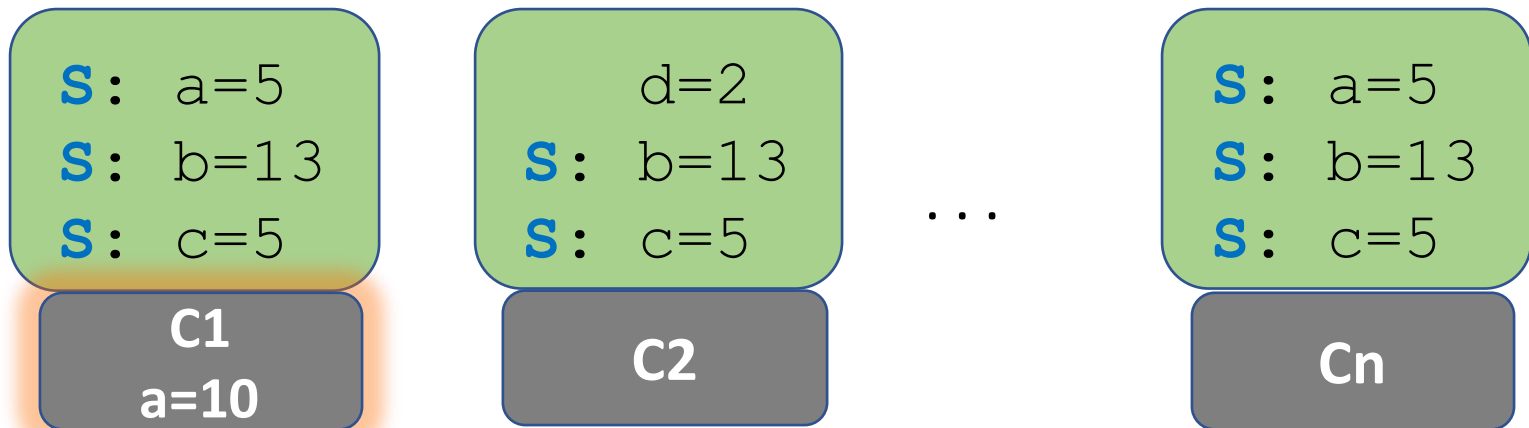


MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

C1 computes new value a=10 in register

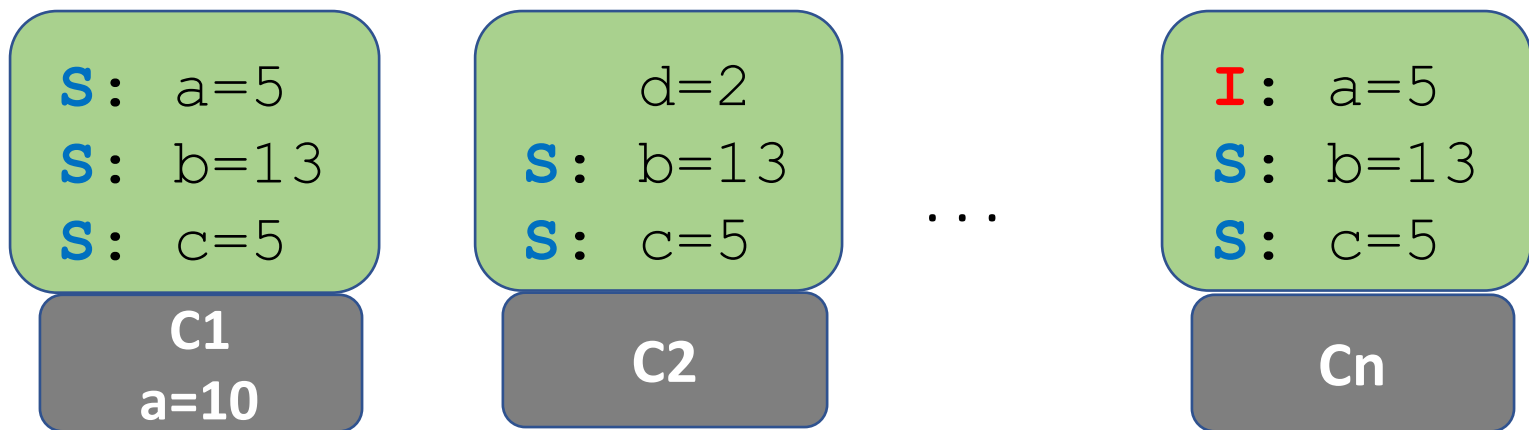


MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Hardware invalidates all other copies of a




MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Hardware updates a in cache of C1 and marks it with 'M'

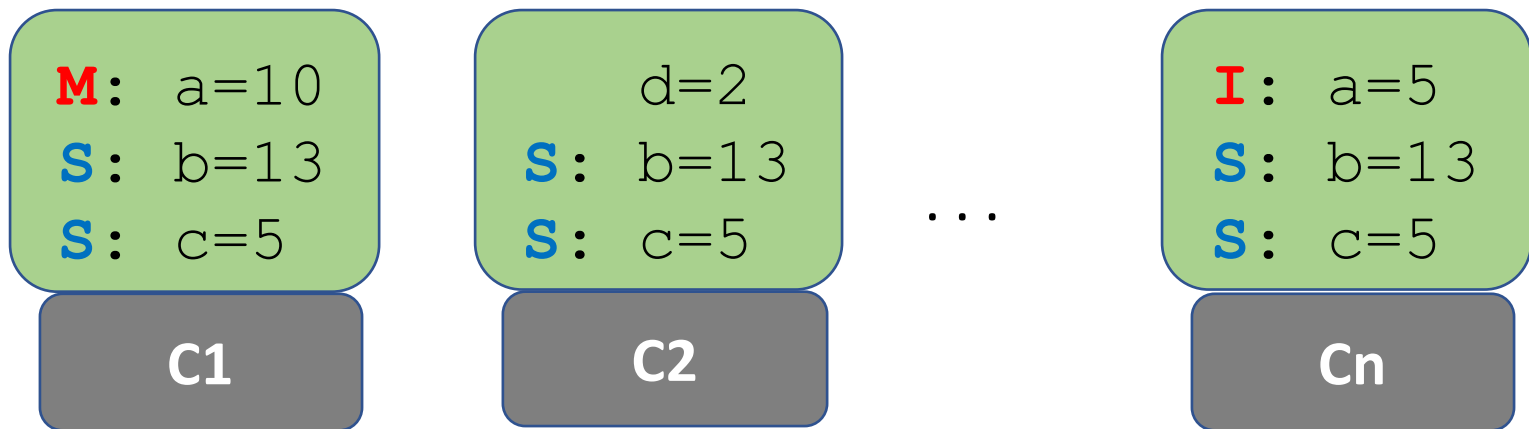


MSI protocols

MSI protocol is a simple cache coherence protocol.
In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

Hardware also updates 'a' in the shared memory



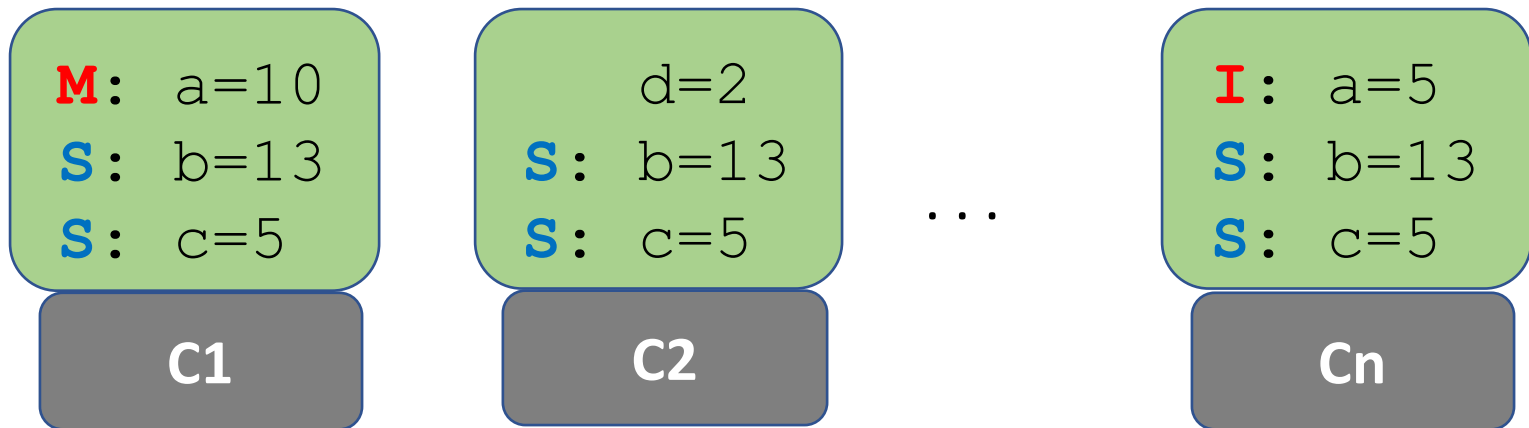
MSI protocols

MSI protocol is a simple cache coherence protocol.

In this protocol, each cache line is labeled with a state:

- **M**: cache block has been modified.
- **S**: other caches may be sharing this block.
- **I**: cache block is invalid

If C_n wants to load 'a' then 'a' will be brought from the shared memory to the local cache of C_n.



Programming multicore platforms

- Option 1: Program directly targeting processor cores
 - Programmer takes care of synchronization
 - Painful and error-prone
- Option 2: Use a *concurrency platform*
 - It abstracts processor cores, handles synchronization and communication protocols, and performs load balancing.
 - Hence offers much easier multicore programming environment
 - Examples:
 - **Pthreads** and WinAPI threads
 - OpenMP